

SQLite

SQLite Manager - /Users/ql...
chrome://sqlitemanager/content/sqlitemanager.xul

Database Table Index View Trigger Tools Help

Directory (Select Profile Database) Go

Structure Browse & Search Execute SQL DB Settings

Enter SQL
SELECT * FROM customer

Run SQL Actions Last Error: not an error

CustomerID	Firstname	Lastname	Sex	Phone
1	Jenny	Ck	F	182990902
2	Kevin	Jo	M	102992883
3	Bob	Rw	M	78279182983
4	Mike	Po	M	8018392121

SQLite 3.14.1 Gecko 51.0.1 0.8.3.1-signed.1-signed Exclusive Number of Rows Returned: 4 ET: 1 ms

SQLite Manager - /Users/ql...
chrome://sqlitemanager/content/sqlitemanager.xul

Database Table Index View Trigger Tools Help

Directory (Select Profile Database) Go

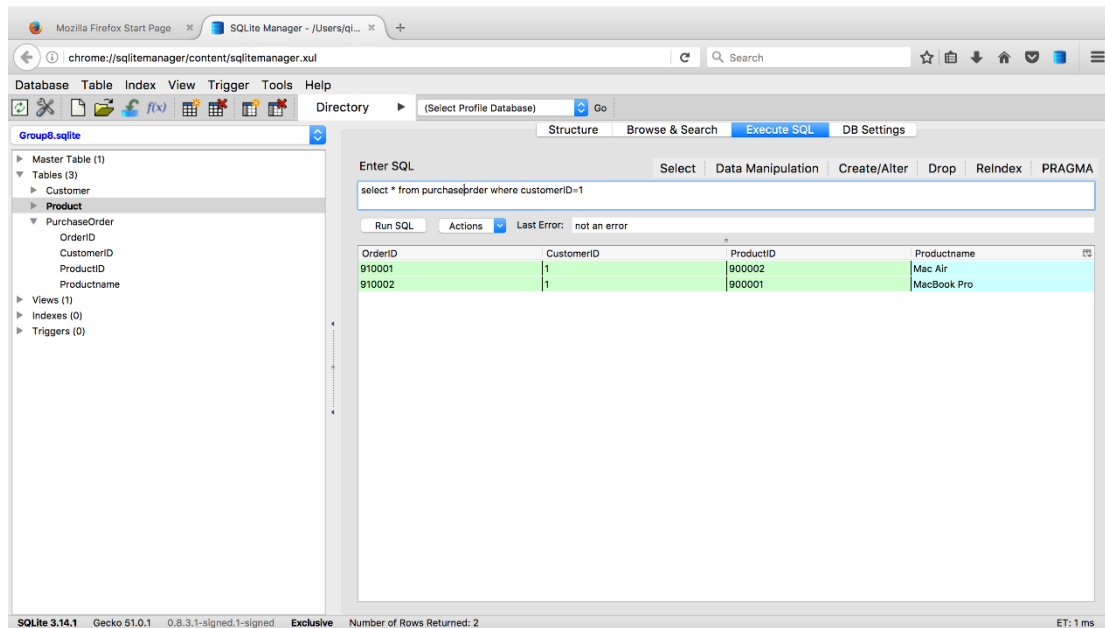
Structure Browse & Search Execute SQL DB Settings

Enter SQL
SELECT * FROM product

Run SQL Actions Last Error: not an error

ProductID	Productname	Price	Size	Color
900001	MacBook Pro17	11,000	15	Grey
900002	Mac Air	10,500	12	light Grey
900003	iPad	800	9	light Grey
900004	iPhone6	12,800	5	black
900005	iPhone7	12,900	5	pink

SQLite 3.14.1 Gecko 51.0.1 0.8.3.1-signed.1-signed Exclusive Number of Rows Returned: 5 ET: 0 ms



In 3rd normal form, Productname should not be included in the PurchaseOrder table as it is already in the product table. However, we chose to include Productname in the PurchaseOrder table just for performance consideration so that when running many queries against the order table, we don't have to join two tables repeatedly.

DB2 Express C

Run a sample query (use where clause and Group by):

```
"BEGIN-OF-STATEMENT". Expected tokens may include: "<call>". SQLS  
db2 => select avg(salary) from emp where edlevel=16 group by sex  
  
1  
-----  
76742.50000000000000000000000000  
54088.50000000000000000000000000  
  
2 record(s) selected.  
  
db2 =>
```

Generate query explain plan (use: db2exfmt tool):

```
db2 => select avg(salary) from emp where edlevel=16 group by sex
SQL0217W The statement was not executed as only Explain information requests
are being processed.  SQLSTATE=01604
db2 => quit
DB20000I The QUIT command completed successfully.
[db2inst1@2167e931c8d5 ~]$ db2exfmt
DB2 Universal Database Version 10.5, 5622-044 (c) Copyright IBM Corp. 1991, 2012
Licensed Material - Program Property of IBM
IBM DATABASE 2 Explain Table Format Tool

Enter Database Name ==> sample
Connecting to the Database.
Connect to Database Successful.
Using SYSTOOLS schema for Explain tables.
Enter up to 26 character Explain timestamp (Default -1) ==>
Enter up to 128 character source name (SOURCE_NAME, Default %) ==>
Enter source schema (SOURCE_SCHEMA, Default %) ==>
Enter section number (0 for all, Default 0) ==>
Enter outfile name. Default is to terminal ==>
DB2 Universal Database Version 10.5, 5622-044 (c) Copyright IBM Corp. 1991, 2012
Licensed Material - Program Property of IBM
IBM DATABASE 2 Explain Table Format Tool

***** EXPLAIN INSTANCE *****

DB2 VERSION:      10.05.5
FORMATTED ON DB:  SAMPLE
SOURCE_NAME:      SQLC2K26
SOURCE_SCHEMA:    NULLID
SOURCE_VERSION:
EXPLAIN_TIME:     2017-02-19-06.23.09.013916
EXPLAIN_REQUESTER: DB2INST1

Database Context:
-----
Parallelism:      None
CPU Speed:        1.889377e-07
Comm Speed:       0
Buffer Pool size: 1000
Sort Heap size:   256
Database Heap size: 1200
Lock List size:   4096
Maximum Lock List: 10
Average Applications: 1
Locks Available:  13107

Package Context:
```

```
Package Context:
-----
SQL Type:         Dynamic
Optimization Level: 5
Blocking:         Block All Cursors
Isolation Level:   Cursor Stability

----- STATEMENT 1 SECTION 201 -----
QUERYNO:          16
QUERYTAG:          CLP
Statement Type:    Select
Updatable:         No
Deletable:         No
Query Degree:      1

Original Statement:
-----
select
avg(salary)
from
emp
where
edlevel=16
group by
sex

Optimized Statement:
-----
SELECT
(Q3.$C0 / Q3.$C1)
FROM
(SELECT
SUM(Q2.SALARY),
COUNT(Q2.SALARY)
FROM
(SELECT
Q1.SEX,
Q1.SALARY
FROM
DB2INST1.EMPLOYEE AS Q1
WHERE
(Q1.EDLEVEL = 16)
) AS Q2
GROUP BY
Q2.SEX
) AS Q3
```

```
Access Plan:
-----
      Total Cost:          6.83611
      Query Degree:       1

      Rows
      RETURN
      ( 1)
      Cost
      I/O
      |
      2
      GRPBY
      ( 2)
      6.83561
      |
      2
      TBSCAN
      ( 3)
      6.83542
      |
      2
      SORT
      ( 4)
      6.83504
      |
      14
      TBSCAN
      ( 5)
      6.83208
      |
      42
      TABLE: DB2INST1
      EMPLOYEE
      Q1

Extended Diagnostic Information:
-----

Diagnostic Identifier: 1
Diagnostic Details:   EXP0073W The following MQT or statistical view was
                        not eligible because one or more data filtering
```

```
Extended Diagnostic Information:
-----

Diagnostic Identifier: 1
Diagnostic Details:   EXP0073W The following MQT or statistical view was
                        not eligible because one or more data filtering
                        predicates from the query could not be matched with
                        the MQT: "DB2INST1"."ADEFUSR".

Diagnostic Identifier: 2
Diagnostic Details:   EXP0148W The following MQT or statistical view was
                        considered in query matching: "DB2INST1"."ADEFUSR".

Plan Details:
-----

      1) RETURN: (Return Result)
      Cumulative Total Cost:          6.83611
      Cumulative CPU Cost:            164654
      Cumulative I/O Cost:             1
      Cumulative Re-Total Cost:        0.0185932
      Cumulative Re-CPU Cost:          98409
      Cumulative Re-I/O Cost:           0
      Cumulative First Row Cost:       6.83556
      Estimated Bufferpool Buffers:     0

      Arguments:
      -----
      BLDLEVEL: (Build level)
      DB2 v10.5.0.5 : s141128
      HEAPUSE : (Maximum Statement Heap Usage)
      112 Pages
      PLANID : (Access plan identifier)
      ea6e1a2e8a8160f3
      PREPTIME: (Statement prepare time)
      59 milliseconds
      SEMEVID : (Semantic environment identifier)
      431f78d03d9bb07e
      STMTHEAP: (Statement heap size)
      8192
      STMTID : (Normalized statement identifier)
      ec6fe7759a569cc0

      Input Streams:
      -----
      5) From Operator #2

      Estimated number of rows:      2
      Number of columns:             1
```

```

-----
5) From Operator #2

      Estimated number of rows:      2
      Number of columns:             1
      Subquery predicate ID:         Not Applicable

      Column Names:
      -----
      +Q4.$C0

2) GRPBY : (Group By)
      Cumulative Total Cost:         6.83561
      Cumulative CPU Cost:           161984
      Cumulative I/O Cost:            1
      Cumulative Re-Total Cost:       0.0188887
      Cumulative Re-CPU Cost:         95739
      Cumulative Re-I/O Cost:         0
      Cumulative First Row Cost:      6.8354
      Estimated Bufferpool Buffers:    0

Arguments:
-----
AGGMODE : (Aggregation Mode)
      FINAL
GROUPBYC: (Group By columns)
      TRUE
GROUPBYN: (Number of Group By columns)
      1
GROUPBYR: (Group By requirement)
      1: Q2.SEX
ONEFETCH: (One Fetch flag)
      FALSE

Input Streams:
-----
4) From Operator #3

      Estimated number of rows:      2
      Number of columns:             2
      Subquery predicate ID:         Not Applicable

      Column Names:
      -----
      +Q2.SEX(A)+Q2.SALARY

Output Streams:

```

```

Output Streams:
-----
5) To Operator #1

      Estimated number of rows:      2
      Number of columns:             1
      Subquery predicate ID:         Not Applicable

      Column Names:
      -----
      +Q4.$C0

3) TBSCAN: (Table Scan)
      Cumulative Total Cost:         6.83542
      Cumulative CPU Cost:           161018
      Cumulative I/O Cost:            1
      Cumulative Re-Total Cost:       0.0179062
      Cumulative Re-CPU Cost:         94773
      Cumulative Re-I/O Cost:         0
      Cumulative First Row Cost:      6.83533
      Estimated Bufferpool Buffers:    0

Arguments:
-----
MAXPAGES: (Maximum pages for prefetch)
      ALL
PREFETCH: (Type of Prefetch)
      NONE
SCANDIR : (Scan Direction)
      FORWARD
SPEED   : (Assumed speed of scan, in sharing structures)
      SLOW
THROTTLE: (Scan may be throttled, for scan sharing)
      FALSE
VISTBLE : (May be included in scan sharing structures)
      FALSE
WRAPPING: (Scan may start anywhere and wrap)
      FALSE

Input Streams:
-----
3) From Operator #4

      Estimated number of rows:      2
      Number of columns:             2
      Subquery predicate ID:         Not Applicable

      Column Names:

```

```

Subquery predicate ID:      Not Applicable

Column Names:
-----
+Q2.SEX(A)+Q2.SALARY

Output Streams:
-----
4) To Operator #2

Estimated number of rows:    2
Number of columns:          2
Subquery predicate ID:      Not Applicable

Column Names:
-----
+Q2.SEX(A)+Q2.SALARY

4) SORT : (Sort)
Cumulative Total Cost:      6.83584
Cumulative CPU Cost:        158981
Cumulative I/O Cost:        1
Cumulative Re-Total Cost:   0.0175213
Cumulative Re-CPU Cost:     92736
Cumulative Re-I/O Cost:     0
Cumulative First Row Cost:  6.83584
Estimated Bufferpool Buffers: 1

Arguments:
-----
AGGMODE : (Aggregation Mode)
PARTIAL
DUPLWARN: (Duplicates Warning flag)
FALSE
NUMROWS : (Estimated number of rows)
2
ROWWIDTH: (Estimated width of rows)
33
SORTKEY : (Sort Key column)
1: Q2.SEX(A)
TEMPSIZE: (Temporary Table Page Size)
8192
UNIQUE  : (Uniqueness required flag)
FALSE

Input Streams:
-----

```

```

FALSE

Input Streams:
-----
2) From Operator #5

Estimated number of rows:    14
Number of columns:          2
Subquery predicate ID:      Not Applicable

Column Names:
-----
+Q2.SALARY+Q2.SEX

Output Streams:
-----
3) To Operator #3

Estimated number of rows:    2
Number of columns:          2
Subquery predicate ID:      Not Applicable

Column Names:
-----
+Q2.SEX(A)+Q2.SALARY

5) TBSCAN: (Table Scan)
Cumulative Total Cost:      6.83208
Cumulative CPU Cost:        143329
Cumulative I/O Cost:        1
Cumulative Re-Total Cost:   0.0175213
Cumulative Re-CPU Cost:     92736
Cumulative Re-I/O Cost:     0
Cumulative First Row Cost:  6.81578
Estimated Bufferpool Buffers: 1

Arguments:
-----
CUR_COMM: (Currently Committed)
TRUE
LCKAVOID: (Lock Avoidance)
TRUE
MAXPAGES: (Maximum pages for prefetch)
ALL
PREFETCH: (Type of Prefetch)
NONE
ROWLOCK : (Row Lock intent)

```

```

Cumulative First Row Cost:      8.81578
Estimated Bufferpool Buffers:   1

Arguments:
-----
CUR_COMM: (Currently Committed)
      TRUE
LCKAVOID: (Lock Avoidance)
      TRUE
MAXPAGES: (Maximum pages for prefetch)
      ALL
PREFETCH: (Type of Prefetch)
      NONE
ROWLOCK : (Row Lock intent)
      SHARE (CS/RS)
SCANDIR  : (Scan Direction)
      FORWARD
SKIP_INS: (Skip Inserted Rows)
      TRUE
SPEED    : (Assumed speed of scan, in sharing structures)
      FAST
TABLOCK  : (Table Lock intent)
      INTENT SHARE
TBISOLVL: (Table access Isolation Level)
      CURSOR STABILITY
THROTTLE: (Scan may be throttled, for scan sharing)
      TRUE
VISIBLE  : (May be included in scan sharing structures)
      TRUE
WRAPPING: (Scan may start anywhere and wrap)
      TRUE

Predicates:
-----
3) Sargable Predicate,
      Comparison Operator:      Equal (=)
      Subquery Input Required:  No
      Filter Factor:           0.333333

      Predicate Text:
      -----
      (Q1.EDLEVEL = 16)

Input Streams:
-----
1) From Object DB2INST1.EMPLOYEE

```

```

Input Streams:
-----
1) From Object DB2INST1.EMPLOYEE

      Estimated number of rows:      42
      Number of columns:             4
      Subquery predicate ID:         Not Applicable

      Column Names:
      -----
      +Q1.$RID$+Q1.SALARY+Q1.SEX+Q1.EDLEVEL

Output Streams:
-----
2) To Operator #4

      Estimated number of rows:      14
      Number of columns:             2
      Subquery predicate ID:         Not Applicable

      Column Names:
      -----
      +Q2.SALARY+Q2.SEX

```

```

Objects Used in Access Plan:
-----

Schema: DB2INST1
Name:   EMP
Type:   Alias (reference only)

Schema: DB2INST1
Name:   ADEFUSR
Type:   Materialized View (reference only)

Schema: DB2INST1
Name:   EMPLOYEE
Type:   Table

      Time of creation:      2017-02-17-03.31.56.395206
      Last statistics update: 2017-02-19-05.42.33.919849
      Number of columns:     14
      Number of rows:        42
      Width of rows:        99
      Number of buffer pool pages: 1
      Number of data partitions: 1
      Distinct row values:   No

```

```

Output Streams:
-----
      2) To Operator #4

      Estimated number of rows:      14
      Number of columns:            2
      Subquery predicate ID:        Not Applicable

      Column Names:
      -----
      +Q2.SALARY+Q2.SEX

Objects Used in Access Plan:
-----

Schema: DB2INST1
Name: EMP
Type: Alias (reference only)

Schema: DB2INST1
Name: ADEFUSR
Type: Materialized View (reference only)

Schema: DB2INST1
Name: EMPLOYEE
Type: Table
      Time of creation:      2017-02-17-03.31.56.399206
      Last statistics update: 2017-02-19-05.42.33.919849
      Number of columns:      14
      Number of rows:         42
      Width of rows:          99
      Number of buffer pool pages: 1
      Number of data partitions: 1
      Distinct row values:     No
      Tablespace name:         USERSPACE1
      Tablespace overhead:     6.725600
      Tablespace transfer rate: 0.080000
      Source for statistics:   Single Node
      Prefetch page count:    32
      Container extent page count: 32
      Table overflow record count: 0
      Table Active Blocks:     -1
      Average Row Compression Ratio: 0
      Percentage Rows Compressed: 0
      Average Compressed Row Size: 0

Executing Connect Reset -- Connect Reset was Successful.
[db2inst1@2167e931cd5 ~]$

```

Graph Data store

```

2. 10.0.0.40 (qhuang)
Re-attach Fullscreen Stay on top Duplicate
qhuang at Qings-Mac in ~/hq/IBM-graph
$ curl -X POST \
  -H "Authorization: gds-token $TOKEN" \
  -H "Content-Type: application/json" \
  -d "$SCHEMA" | jq '.'
{
  "requestId": "21ff6099-1b27-4a0f-823b-320fbf14eed0",
  "status": {
    "message": "",
    "code": 200,
    "attributes": {}
  },
  "result": {
    "data": [
      {
        "propertyKeys": [
          {
            "name": "name",
            "dataType": "String",
            "cardinality": "SINGLE"
          },
          {
            "name": "verified",
            "dataType": "Boolean",
            "cardinality": "SINGLE"
          },
          {
            "name": "tweet",
            "dataType": "String",
            "cardinality": "SINGLE"
          },
          {
            "name": "sentiment",
            "dataType": "String",
            "cardinality": "SINGLE"
          },
          {
            "name": "tone",
            "dataType": "String",
            "cardinality": "SINGLE"
          },
          {
            "name": "hashtag",
            "dataType": "String",
            "cardinality": "SINGLE"
          },
          {
            "name": "numTimes",
            "dataType": "Integer",
            "cardinality": "SINGLE"
          },
          {
            "name": "time",
            "dataType": "String",
            "cardinality": "SINGLE"
          }
        ],
        "vertexLabels": [
          {
            "name": "person"
          }
        ]
      }
    ]
  }
}

```



```
2. 10.0.0.40 (qhuang)
Re-attach Fullscreen Stay on top Duplicate
{
  "name": "person"
},
{
  "name": "hashtag"
},
{
  "name": "tweet"
}
],
"edgeLabels": [
{
  "name": "mentions",
  "directed": true,
  "multiplicity": "MULTI"
},
{
  "name": "hashes",
  "directed": true,
  "multiplicity": "MULTI"
},
{
  "name": "tweets",
  "directed": true,
  "multiplicity": "MULTI"
},
{
  "name": "favorites",
  "directed": true,
  "multiplicity": "MULTI"
}
],
"vertexIndexes": [
{
  "name": "vByName",
  "composite": true,
  "unique": true,
  "propertyKeys": [
    "name"
  ],
  "requiresReindex": false,
  "type": "vertex"
},
{
  "name": "vByVerified",
  "composite": true,
  "unique": false,
  "propertyKeys": [
    "verified"
  ],
  "requiresReindex": false,
  "type": "vertex"
},
{
  "name": "vBySentiment",
  "composite": true,
  "unique": false,
  "propertyKeys": [
    "sentiment"
  ],
  "requiresReindex": false,
  "type": "vertex"
}
],
{
```

```
2. 10.0.0.40 (qhuang)
Re-attach Fullscreen Stay on top Duplicate
"requiresReindex": false,
"type": "vertex"
},
{
  "name": "vByTone",
  "composite": true,
  "unique": false,
  "propertyKeys": [
    "tone"
  ],
  "requiresReindex": false,
  "type": "vertex"
},
{
  "name": "vByTweet",
  "composite": true,
  "unique": false,
  "propertyKeys": [
    "tweet"
  ],
  "requiresReindex": false,
  "type": "vertex"
},
{
  "name": "vByHashtag",
  "composite": true,
  "unique": true,
  "propertyKeys": [
    "hashtag"
  ],
  "requiresReindex": false,
  "type": "vertex"
},
{
  "name": "vByNumTimes",
  "composite": true,
  "unique": false,
  "propertyKeys": [
    "numTimes"
  ],
  "requiresReindex": false,
  "type": "vertex"
}
],
"edgeIndexes": [
  {
    "name": "eByTime",
    "composite": true,
    "unique": false,
    "propertyKeys": [
      "time"
    ],
    "requiresReindex": false,
    "type": "edge"
  }
]
},
"meta": {}
}
}

qhuang at Qinas-Mac in ~/hg/IBM-graph
$ cat << ENDGREMLIN >gremlin.json # write everything until ENDGREMLIN into gremlin.json
```

```
2. 10.0.0.40 (qhuang)
Re-attach Fullscreen Stay on top Duplicate
def prachi = graph.addVertex(T.label, 'person', 'name', 'Prachi', 'verified', false);
iasm, 'tone', 'excited');h.addVertex(T.label, 'tweet', 'tweet', 'I adore soccer #exercise', 'sentiment', 'enthus
def dancingHashtag = graph.addVertex(T.label, 'hashtag', 'hashtag', 'dancing', 'numTimes', 2187);
def exerciseHashtag = graph.addVertex(T.label, 'hashtag', 'hashtag', 'exercise', 'numTimes', 7);
def keith = graph.addVertex(T.label, 'person', 'name', 'Keith', 'verified', false);
n', 'tone', 'declarative');ddVertex(T.label, 'tweet', 'tweet', '@Matt Beer is #delicious', 'sentiment', 'adoratio
def matt = graph.addVertex(T.label, 'person', 'name', 'Matt', 'verified', false);
def deliciousHashtag = graph.addVertex(T.label, 'hashtag', 'hashtag', 'delicious', 'numTimes', 293);
david.addEdge('tweets', b- david.addEdge('tweets', browniesTweet);
browniesTweet.addEdge('hashes', bakingHashtag);
browniesTweet.a- browniesTweet.addEdge('mentions', joseph);
joseph.addEdge('favorites', browniesTweet, 'time', '10:30PM EST');
joseph.addEdge('tweets', graphTweet);
graphTweet.addEdge('hashes', blueMixHashtag);
kim.addEd- kim.addEdge('favorites', graphTweet, 'time', '4:29AM EST');
kim.addEdge('favorites', graphTweet, 'time', '12:21AM EST');
kamal.add- kamal.addEdge('favorites', dancingTweet, 'time', '9:17PM EST');
kamal.addEdge('tweets', soccerTweet);
prachi.addEdge('tweets', dancingTweet);
dancingTweet.addEdge('hashes', dancingHashtag);
dancingTwe- dancingTweet.addEdge('hashes', exerciseHashtag);
soccerTweet.addEdge('hashes', exerciseHashtag);
keith.addEdge('tweets', beerTweet);
beerTweet.addEdge('mentions', matt);
beerTweet.addEdge('hashes', deliciousHashtag);
kamal.addEdge('favorites', graphTweet, 'time', '5:22PM UTC');
}
}
ENDGREMLIN

qhuang at Qings-Mac in ~/hq/IBM-graph
$ curl -X POST -H "Authorization: gds-token $TOKEN" -H "Content-Type: application/json" -d @gremlin.json | jq '.'
{
  "requestId": "ec265c66-eaac-40f4-93b9-bcfe32924064",
  "status": {
    "message": "",
    "code": 200,
    "attributes": {}
  },
  "result": {
    "data": [
      {
        "id": "bgj-cns-bv9-39c",
        "label": "favorites",
        "type": "edge",
        "inVLabel": "tweet",
        "outVLabel": "person",
        "inV": 4224,
        "outV": 16408,
        "properties": {
          "time": "5:22PM UTC"
        }
      }
    ],
    "meta": {}
  }
}

qhuang at Qings-Mac in ~/hq/IBM-graph
$
```

IBM Graph

GraphDB > test

graph test

```
graph.traversal().V().hasLabel('person').has('name', 'Kamal').out('favorites').values().path();
```

Filter: Label Type Properties

Vertices: 3

graph test

```
graph.traversal().V().hasLabel('person').has('name', 'Kamal').outE('favorites').values().path();
```

GraphDB > test

1 // Enter your Gremlin query here. (Shift + Enter) to execute.

Graph test

```
graph.traversal().V().hasLabel('person').has('name', 'kamel').outE().inV();
```

63 +
64 {
65 "id": "Bar-fik-3yd",
66 "value": "happy"
67 }
68 }
69 +
70 {
71 "id": "7ib-fik-3dh",
72 "value": "I love ballroom dancing & exercise"
73 }
74 }
75 }
76 }
77 }

Filter Label Type Properties

Vertices: 2

Graph test

```
graph.traversal().V().hasLabel('person').has('name', 'kamel');
```

7 +
8 {
9 "id": "5af-cns-1l",
10 "value": "kamel"
11 }
12 }
13 +
14 {
15 "id": "5bn-cns-1ll",
16 "value": "kamel"
17 }

Filter Label Type Properties

Vertices: 1

GraphDB > test

1 // Enter your Gremlin query here. (Shift + Enter) to execute.

Graph test

```
graph.traversal().V().hasLabel('person').has('name', 'kamel').outE('favorites').inV();
```

1 +
2 {
3 "id": 4234,
4 "label": "twent",
5 "type": "vertex",
6 "properties": {
7 "sentiment": {
8 {
9 "id": "1ic-39c-3ts",
10 "value": "enthusiasm"
11 }
12 }
13 }
14 "name": {
15 "id": "1ik-39c-3yd",
16 "value": "kamel"
17 }
18 }
19 }
20 }
21 }

Filter Label Type Properties

Vertices: 2